

Automate the Intake. Preserve the Thinking.

A Philosophy for Working With AI Without Losing Yourself.

Graham Friedman

Abstract. This paper is a personal account of what it actually takes to work with AI well, written from inside the first recruiting cycle meaningfully affected by AI at scale. It argues that context architecture is the core skill in AI-assisted knowledge work, not prompt engineering alone, although it does matter. It comments on the philosophy of automation workflows and traces the origin of a three-tier framework called Automate, Amplify, Preserve. This framework was arrived at independently through building and failing, and shows why that framework is as relevant to how companies deploy AI systems to face their clients as it is to how individuals use AI day to day trying to improve the efficiency of their work. It argues that as AI floods every channel with generic output, the human elements we are most at risk of automating away are also the ones that matter most: voice, judgment, and real relationships.

Contents

- I. The Problem Nobody Is Naming
- II. A Framework That Already Existed
- III. Context Is the Skill, And So Is Prompting
- IV. What AI Cannot Do
- V. The Framework in Practice: What Companies Are Getting Wrong
- VI. The Broader Ripple: The Bubble Debate, Physical Limits, and What Everyone Is Missing
- VII. What AI Cannot Touch: Relationships, Voice, and the Disconnection Crisis
- VIII. A Personal Philosophy for the AI Era

I. The Problem Nobody Is Naming

Being in the AI space every day, I turn my phone on and immediately see a myriad of varying perspectives and news. What strikes me is not any single take. It is how loud and how contradictory all of it is at once. Open LinkedIn and scroll for a minute, especially in the tech space. Someone explains what AI is as if nobody has heard of it. Someone else predicts total economic collapse. Someone sells a course that promises to triple your income with automation. Some people point to the fact that AI really is not helping them profit more, it is simply increasing the volume they can produce. Right next to all of that, on a more personal level, senior workers are either lamenting that younger colleagues cannot do real work anymore, or they are praising how that same generation makes the busywork and repetitive tasks disappear. There are so many experts, so much fear of being left behind, and so much opposition to adopting AI. And there are so many people who understand it all just as minimally as the rest, but who are adamant in their opinions. All of it, every day, all at one time, in one feed. That is the discourse.

The polarization is not the interesting part. What gets me is how little of it is grounded in what people are actually doing with these tools. Most of the fear is performative, and a lot of it is also just promotion. The warning and the pitch are often in the same post. Meanwhile, the people actually in the middle, building specific things and testing them in real work with real implications, are the ones who understand the technology most clearly, both its benefits and its shortcomings. And they are almost invisible, even though that is where the only interesting work is happening.

It is worth being honest about why the noise is this loud, because a lot of it is not accidental. Both the fear and the hype serve someone financially. That is not a conspiracy, it is just how incentives work. AI companies benefit from urgency. The more people believe they are falling behind, the more subscriptions get bought, the more enterprise contracts get signed, the more courses get sold, the more funding rounds get justified. Fear sells better than a measured take. When the CEO of a major lab publishes an essay warning that AI could wipe out half of entry-level white-collar jobs, that is not neutral information. He might be right. But he is also the person selling the product, and reading it as straight analysis instead of interested communication is a mistake. The same thing happens on the other end. A researcher who warns of existential risk gets more attention and more funding than one who says it depends on the use case. "AI will replace everyone" and "AI is all hype" both get engagement. Nuance gets ignored.

So the loudest voices get rewarded for being loud, and the careful work mostly stays quiet. The only way through that is to do your own work. Run the systems yourself. Test the claims against what you actually see happen. Form your own read instead of borrowing someone else's, especially when that someone has something to sell you. It is harder than scrolling AI content to simply feel productive and valuable. It is also the only way to actually understand what is going on.

And here is the part that gets lost: most people are not nearly as far behind as they feel. The scale is real. ChatGPT alone crossed roughly 900 million weekly active users by early 2026, more than double a year earlier, and close to a billion people now use AI chatbots worldwide. (OpenAI via TechCrunch, 2026; DemandSage, 2026) But scale is not the same as depth. When you look at what people actually do with these tools, the top use across US adults is searching for information, largely assisting consumerism, around 60%, followed by generating ideas, with coding, automation, and system building far down the list. (Pew, 2026) In the UK, the single most common reason people give for using generative AI is that it is fun, ahead of learning or research. (Ofcom, 2026) The vast majority of what gets called "AI use" worldwide is really a more convenient, more conversational version of Google, or a judgment-free companion to talk to.

It is not just teenagers, and it is not evenly spread. Adoption tracks age almost linearly. Deloitte found 76% of Gen Z have used standalone generative AI tools, against 58% of millennials, 36% of Gen X, and 20% of boomers. (Deloitte, 2026) Interestingly, the most frequent daily users are millennials, not Gen Z, at 37% daily versus 25%. (Edubrain via Fast Company, 2026) Gen Z leans on it most for school, with 61% using it for learning, while Gen X and boomers reach for it mostly for work efficiency and basic convenience. (SurveyMonkey, 2026) Across every one of those groups, the dominant behavior is the same: ask a question, get an answer, move on. Very few people are building systems, designing context, or thinking hard about where the tool should and should not be trusted. The real work is rarer than the usage numbers suggest. If you are reading this and feel behind, you are almost certainly further along than you think.

I came back from studying abroad in Japan late into the first recruiting cycle that was meaningfully affected by AI at scale. Every major deadline had passed. And I was figuring out, for the first time, how structured recruiting actually works, at exactly the moment it stopped working the way it had for the previous decade. I was also far too late for the companies I was initially targeting. The tools that promised to help me were everywhere. Mass-apply extensions, AI resume builders, cover letter generators. The pitch was always the same: efficiency at scale. Apply to 200 companies instead of 20. So I watched what happened when people used them. Structurally correct outputs. Completely generic. Not aligned to who the person actually was, or what they wanted to do, or what company and culture fit them the best. It was simply a way to land an interview at the first company that gave your application, the one you only glanced at once, a chance. The Greenhouse CEO describes a "doom loop" where candidates use more AI to break through filters, employers deploy stronger filters, and candidates respond with even more automation. (Shortlistd, 2026) And no one ends up satisfied. The whole workforce is misaligned and potentially less effective because of it. One 2026 hiring survey found 80% of hiring managers can identify AI-written resumes at a glance. The figure below captures the dynamic: application volume climbing year over year while the quality of signal reaching recruiters collapses in the opposite direction.

Here is my opinion in short:

"These tools are not just unhelpful. They are a major contributor to the noise recruiting teams are still trying to manage, and the main contributor to the noise they promise to help you fight through."




Figure 1. Application volume vs. recruiter signal quality, indexed to 2022. Synthesized from Metaintro (2026), CoderPad State of Tech Hiring (2026), Greenhouse research.

300+ avg applications per job posting in 2026	80% of hiring managers who spot AI resumes at a glance	900M weekly ChatGPT users, but mostly basic search and chat
--	---	--

Sources: Metaintro (2026); 2026 recruiting survey; OpenAI via TechCrunch (2026).

The response I kept coming back to was not to just apply to more places. It was to figure out how to build something that could produce outputs that could only have come from me. Speed on the mechanical work, preserved judgment on everything requiring my actual voice, my real story, and my genuine fit with the people I was reaching out to. All of this while still understanding that without internal connections and advocates, volume does matter to some extent. Building a system that balanced the two perfectly together, and one that expands on the context as I gain experience so I can use it forever.

As expected, being on the internet, there is a lot of complaining about the current system, and a lot of this is justified and grounded, despite being heavily exaggerated for comedic purposes. There are endless videos of parents who cannot understand why their kids cannot get a job, giving advice like just walk into the building and ask, they will be impressed, or give a strong handshake, talk about your experience, tell them you would really really love to work there, send a thank you note. None of that is wrong. It is just the bare minimum now. They could walk in, have someone glance at a resume, and in some more entry level cases get hired on a related background and some enthusiasm. There was also a much smaller pool of competition for almost

every role. And there was much, much less of having to prove they already knew how to do the job before they got it. There used to be more resources and systems put in place to train people. Now, especially with AI, companies do not want to spend the money or time to do that. So candidates are forced to become employable completely independently, and not everyone is cut out to operate outside of a system and be fully self motivated. Not to mention, especially in tech, you can often find yourself live coding in front of people who design questions specifically for you to fail, and you are expected to have every skill and prove you can do the job before anyone lets you do it. It hits entry level the hardest. Training people takes time, and companies are realizing they can just hand those entry-level tasks to AI, so intern classes are smaller and youth underemployment is higher than ever. The seniors and the parents giving the advice are often far removed from what getting a job actually looks like today. It is so much harder than it was for them. This is why I also see people in senior positions really dislike their job. But since they hear so much about people taking months to a year finding a new position, and not always a better one, they decide to stay in it, because they are afraid they are not cut out for the current state of recruiting even this late into their career.

II. A Framework That Already Existed

I built a system. Then I watched it fail. Then I rebuilt it from scratch. Somewhere in that I landed on a consistent logic for where AI belongs and where it does not, and I called it Automate, Amplify, Preserve. Some months later I came across a practitioner white paper from Rob Murray and Derek Kopen at SapientNitro, published in 2016. Their framework was identical. They built it studying the hotel concierge in the age of digital marketing. I built it debugging a failed job search system in the age of large language models. Same three tiers, same underlying logic, different decade, different context.

When I found their paper I felt good about it, not deflated. Two people working a completely different problem, nine years earlier, in digital marketing instead of LLMs, landed on the exact same three tiers I did. That is not a coincidence, and it is not a clever way to organize a slide. It is describing something real about how humans and machines should split work. Some tasks are mechanical enough that automation is safe. Some need the machine and the human together. And some have a human element so essential that the technology should not touch the moment at all. That maps onto an actual constraint, which is why it keeps holding up.

On top of that system, I added the distinction between Technology Leading and Human Leading because the three tiers alone never told me how to actually run a task. Knowing something belongs in Amplify does not tell you who starts. That ended up being the most useful part of the whole thing in practice, and it is the part I lean on most when I think about how a company should design its own systems.

AUTOMATE	AMPLIFY	PRESERVE
Rule-based and mechanical. Claude processes inputs and produces structured outputs. No personal voice or human judgment required. The output could have been produced by any version of this system for any candidate. Examples: tracker updates, portal formatting, basic research.	Technology and human together. Technology Leading: Claude does the heavy lifting, drafting structured output following rules and frameworks, and the human reviews and makes light adjustments where it does not quite land. Human Leading: the human brings the voice, reasoning, or raw unpolished thinking first, and Claude structures and refines what the human brings.	Human element so essential technology should not touch the moment itself. Relationship, trust, empathy, and real-time judgment cannot be replicated. Claude assists only in preparation before and debrief after.

Table 1. The Automate, Amplify, Preserve framework. Murray & Kopen (SapientNitro, 2016) and independently, this paper.




Figure 2. Human involvement by tier. Automate operates near zero; Preserve is nearly fully human.

The figure above shows the rough split, but the number that matters most is in the middle. The critical distinction within Amplify is whether technology is doing the work and the human is checking it (Technology Leading) or the human is doing the work and technology is structuring it (Human Leading). Resume tailoring is Technology Leading. The rules are clear enough that the system can draft and I review where the voice is off. A prompt there looks like: "Using my full background and the attached job description, tailor my experience section. Follow the writing rulebook exactly. Flag anything you are unsure represents me accurately." The AI leads, I correct.

Writing an outreach message to a specific person is Human Leading. The personal connection, the specific reference, the genuine ask matched to their seniority has to come from me first. There the prompt looks

completely different: "Here is what I actually want to say to this person, in my own rough words, and here is the specific context of how we connected. Do not add anything I did not say. Structure it, tighten it, and tell me if anything reads as generic." I lead, the AI refines. Same tier, opposite starting point. Getting that backwards is what produces outreach that sounds like a template, and people can feel a template instantly. This is the part of the framework I leaned on most heavily in the showcase I built for Prophet's AI Foundry team, because it is the part that translates directly to how a company decides which steps of a client workflow the system should lead and which a human must.

The SapientNitro paper warned in 2016 about three mistakes brands make when replicating human service with technology: playing lip service to personalization without delivering it, treating personalization as exclusivity rather than quality, and over-humanizing digital interactions in ways that trigger the uncanny valley effect. All three mistakes look exactly the same in 2026 AI deployments. The paper warned about this. Most people ignored it.

III. Context Is the Skill, And So Is Prompting

The most common AI advice right now is some version of "learn to prompt better," or blanket statements about here is how you connect and automate everything step by step, take my course or follow me to learn how. But there is so much nuance that goes completely missing. Their advice is not entirely wrong. It just misidentifies where most of the leverage is, and sells you a simpler idea than the reality. Andrej Karpathy, one of the most respected practitioners working in LLMs, describes the knowledge inside a neural network as a vague recollection of things the model saw during training, like something you read a month ago, and the context window as working memory, the stuff directly available to you right now. The implication: what you put in the context window is not just a prompt. It is what the model actually has access to. Generic context produces generic output not because the prompt was vague but because the system did not know enough.

A research paper published in April 2026 made the same argument formally. Context engineering, not prompt engineering, is the core skill in serious LLM applications. The empirical foundation was 200 documented interactions all showing the same pattern: failures were caused by incomplete context, not by poorly constructed prompts. (Context Engineering: A Practitioner Methodology, arXiv 2026)

"The iterations are not caused by poor prompting. They appear to be caused by incomplete context." Context Engineering: A Practitioner Methodology for Structured Human-AI Collaboration, arXiv, April 2026

So I built the system around that principle, and the detail is the point. It is a Claude Project with eight permanent files loaded simultaneously in every session. The first is my full base resume as a reference anchor. The second is a master background and pre-screening document holding exact technical depth on all

six of my projects, the methods used, the sample sizes, the specific findings, written so the system can pull real specifics instead of inventing plausible ones. The third is a complete pipeline file tracking every interaction I have had with every company in the search, who I spoke to, what was said, what stage each is at, so outreach is never blind to history. The fourth is a full writing rulebook governing voice. The fifth is a detailed first-person account of my actual AI experience and reasoning, including where I have watched it fail. The remaining files cover my LinkedIn presence, my GitHub, and a running document of everything I have learned about the job process itself. The whole person, not a resume summary. And I cannot stress this enough: it holds every detail of every step, down to the emotions and the exact thinking behind each decision, in each situation, and after each new conclusion. On top of that sits a JD Analyzer script I built in Claude Code that calls the Anthropic API directly with that same background as a system prompt and returns a structured five-section brief, top skills, core themes, best projects to lead with, recommended framing angle, and red flags, in under sixty seconds per application.

My read on where to draw these lines did not come from theory. Before this, I worked at Data Annotation, where the actual job was reviewing and rating AI outputs, writing better responses than the model produced, and catching the specific places where confident-sounding answers were subtly or completely wrong. Months of that teaches you something no amount of reading about LLMs can: exactly where these systems break, how plausible a wrong answer can look, and how often the failure is not the model being dumb but the model not having what it needed to be right. That experience is the entire foundation of this philosophy. I did not theorize that context is the constraint. I watched it be the constraint, thousands of times, and got paid to fix it.

There are also rules for how the context itself gets formed, because a pile of documents is not architecture. Everything routes through a single source of truth in my notes, then gets distilled into the specific file it belongs in, then gets repasted into the Project so the live system always reflects the latest state. Files are written in the voice and structure the system needs to consume, not the way I would casually jot something down, but still built on my rough thoughts and language. When I learn something durable about the process, it goes into the learnings file immediately rather than living in my head. The architecture is maintained, not just assembled once, and that maintenance is most of why it keeps producing things that actually sound like me.

The way I actually pull that context out of my own head matters as much as the files themselves. I do not start by telling Claude what to write. I start by uploading everything, the raw material all at once: resume, project files, course history, work documentation, nothing held back. Then I ask Claude to ask me, to generate every question it needs to fully understand each piece before it produces anything. The question sets it comes back with are extensive, covering the thought processes, the feelings during specific moments, the technical choices and the reasoning behind them, what each environment actually felt like emotionally and physically, what surprised me, what frustrated me, what felt satisfying, and what a reader would need to understand to use the information correctly. The goal is never just the facts. It is the full human experience behind every decision. I answer all of it in my own unpolished voice, because the raw thinking is the point, not a clean summary. Then I tell Claude to preserve as much of my actual language and structure as possible while

making the output clean and specific, refinement rather than replacement. I proofread and push back on anything that does not sound right or is missing nuance, and I repeat that until every section is accurate and fully defensible. Through all of it I am deliberately redundant, reminding Claude of the structural rules and being explicit in every ask, because specificity is the entire skill. Generic inputs produce generic outputs.

A quick word on the tooling, because the choice matters more than people think. The JD Analyzer runs in Claude Code, calling the API directly, and that is the right call for what it does: a repeatable script I run dozens of times against the same rich background, where I want raw speed and the same structured output every time. This output is also not submitted, it is read for insight, it is not supposed to reflect my actual person. Other processes like this are just kept and added to a mechanical system loop. I would not build the whole system that way, and I think a lot of people reach for the agentic, code-it-all-up approach when they should not. The parts of my system that actually need judgment, the resume tailoring, the outreach, the parts where I am reading every line and pushing back, live in a conversational Project on purpose. A fully automated Claude Code pipeline is great when the task is well-defined and you trust the output enough to not read all of it. The moment the task is judgment-heavy, or the cost of a generic output is high, the automation that felt efficient becomes the thing that strips your voice out. So the tool follows the tier. Automate steps can run as code. Amplify and Preserve steps stay in the loop with me. Picking the heaviest automation available for every step is one of the most common ways these systems quietly go generic, and break when new updates come out, or a new task to solve is requested.

That said, prompting matters too. A vague prompt on top of a rich context still produces a vague output. Being very specific about what you need and why, asking the AI to ask you questions before it produces anything, being deliberately redundant on rules, providing explicit formats and constraints, all of that compounds with context to produce outputs that are genuinely better rather than just faster. The single highest-leverage move is making the model ask you questions first, because it surfaces the gaps in your own thinking before they get baked into a confident, wrong answer. Giving it a worked example of what good looks like beats describing good in the abstract. Telling it explicitly what not to do is often as valuable as telling it what to do. And asking it to reason step by step before it commits to an answer keeps it from cramming everything into one rushed leap. None of these are tricks. They are just ways of being clear, and clarity is what most prompts are missing. Cost is also real. Per-token API pricing means inefficient prompting is literally more expensive. A poorly structured prompt that takes three iterations to correct costs significantly more than one that lands on the first try. (Anthropic API, 2026: \$5 per million input tokens, \$25 per million output tokens.)

The writing rulebook file that tries to replicate my writing illustrates the architectural point precisely. There is a lot of caution against over-humanizing, and I do not disagree. I just argue that the gap between AI output and human voice is not a prompting problem. You can tell AI to "write like me" in every prompt and get outputs that are cleaner, more structured, and completely unlike you. The rulebook I use covers sentence-level patterns specific to how I actually write, a comprehensive banned phrase list (leveraging, utilizing, seamless,

passionate about, proven track record, all of it), detailed output rules for six specific writing types, and a 12-question self-check before any output goes anywhere. What the rulebook identified: I write long sentences accumulating detail through parentheticals, then land short. I think while writing, not before. The outputs still do not match that exactly, but they are much closer with the rulebook than without it. And the gap that remains is exactly where the actual writing work happens. The rulebook does not replace rereading and rewriting. It narrows the starting gap so I can save costs and time, and write more cleanly on the first go. I still have to read every output and push back when it sounds generic or off. Always.

There is one more piece of this that follows directly from the working-memory idea, and it is the thing most people skip: giving the model actual tools. Karpathy is blunt about where these systems fail. They are bad at counting, bad at spelling, bad at mental arithmetic, because each token only gets a finite slice of computation and they cannot see individual characters, only tokens. But the moment you let the model reach for a tool, a web search to pull live information into the context window, or a code interpreter to actually run the calculation instead of guessing it, the failure mode mostly disappears. You are not asking it to remember. You are giving it working memory and a calculator. In my own system this is why the JD Analyzer is a script and not a vibe: the structured, repeatable parts get handed to code, and the model is asked to reason over real pasted-in context rather than recall it. Adding the right tools and skills to an LLM is often a bigger jump in quality than any prompt tweak. But none of that removes the last step. A tool can fetch the right number and the model can still frame it wrong, so you check it. Tools reduce the failures you can predict. They do not replace the human reading the output at the end.




Figure 3. AI business adoption vs. measured productivity gains. The productivity paradox, technology visible everywhere but not yet in the statistics, first identified by Solow in 1987 about computers. Duke University / Federal Reserve CFO Survey (2026).

The productivity numbers are worth sitting with, and the figure above lays out the gap directly: adoption has climbed toward near-universal while the productivity gains companies actually report stay strikingly small. A

2026 Duke and Federal Reserve working paper surveyed 750 CFOs: companies reporting AI-driven productivity gains averaging 1.8%, but revenue data shows much smaller real gains across every major industry. Oxford Economics found firms are not replacing workers with AI on any significant scale and suspects some are using the framing to dress up layoffs. Cost is the other half of the story. Uber burned through its entire 2026 AI budget in four months as Claude Code adoption jumped from 32% to 84% of its 5,000-engineer organization, at \$500 to \$2,000 per engineer per month. (Investing.com, 2026) Some companies that aggressively replaced workers with AI are quietly starting to rehire, and Forrester expects half of AI-attributed layoffs to be reversed. The economics of substitution turned out messier than the headlines suggested.

IV. What AI Cannot Do

The first version of my system failed in a specific and instructive way. Seven automated steps calling the Anthropic API directly. Resume tailoring, portal answers, networking outreach, pipeline tracking, all running sequentially. The outputs were structurally correct but faulty and generic. And when the system broke, you now had to work within the confines of the broken system, you could not work alongside it. It was the same thin-prefix problem from before, just automated end to end: the system did not know the person it was supposed to represent. Rebuilding around a Claude Project copy-and-paste system holding all eight files simultaneously solved everything at once.

What can AI genuinely not do, even with full context? It cannot originate the authentic human reasoning that makes an output recognizably yours. It cannot calibrate to the specific relational context of a single person it has never interacted with. It cannot replicate real-time judgment from actually being in a room with someone. And it will fail on specific things in ways that are surprising. Karpathy documents this as "jagged intelligence": models that solve PhD-level physics problems and then fail to correctly compare 9.11 and 9.9, because those numbers pattern-match to Bible verse notation internally. This is exactly the kind of failure I spent months catching at Data Annotation. The answer looks completely authoritative, the formatting is clean, the tone is confident, and it is simply wrong in a way you only catch if you already know the right answer or check carefully. That is the whole danger. You have to check the work. Every time. Especially when the output looks polished and confident, because polish is exactly what these systems are best at and correctness is what they are least reliable at.

The ownership question matters here, and it opens up a set of questions I do not think we have good answers to yet. Can you take credit for something AI helped produce? At what point does someone get to ask what your abilities actually are? If the most valuable skill is now just being the best at delegating your thinking to a machine, is that a skill we should even be trying to build? Are people quietly working toward doing less and less of their own thinking, and getting worse at it in the process? If you work in AI, are you building your own replacement, or everyone's? I do not think these have clean answers, and there is no real system in

place to settle any of them. AI detectors are wildly inconsistent. There is no agreed-on way to assign ownership. My own read is that you can reasonably claim work as yours if you supplied the reasoning, the voice, the direction, and the judgment, because coming up with what to ask for is itself a form of thinking. For me specifically, the entire system runs on my background and I am writing my voice into every output from scratch, so the claim feels strong. The claim goes hollow the moment someone takes whatever the tool generated without reading it or asking whether it represents them. That is exactly what floods LinkedIn and Instagram right now. Polished, empty content. Courses built on AI slides that cannot deliver what they sell. Expertise that collapses under one real follow-up question. AI makes it possible to claim almost anything and back it up with something that sounds right, which is why sales, selling the image, has quietly become the skill people optimize for instead of the underlying work. And when everyone reaches for the same generic prompts, the outputs converge on the same generic average, so standing out now takes more specific, intentional, human input than it ever did before, not less.

There is a study that makes this concrete. Russell and colleagues compared 61,608 human stories against 5 AI models, and even after stripping out vocabulary and surface style, AI writing was identifiable 93% of the time, purely from how the writing was built. (StoryScope, arXiv 2026) The tells are structural, not lexical. AI spells out the moral instead of letting the reader infer it, 77% of the time against 52% for humans. It describes emotion as body sensations, tight chest, pounding heart, rather than just naming the feeling, 81% against 38%. It keeps one clean storyline and resolves it neatly, has the hero drive the ending through deliberate action, wraps everything up calm and settled. Humans talk straight to the reader, say "I was scared" without dressing it up, name real brands and real places, write people who are morally messy, and leave a second thread half-finished. What that study is really measuring is the gap between producing something competent and producing something that could only have come from one specific person. The competent version is now free and infinite. The specific version is the only thing left that is scarce. It is a shame, since amazing writers who have worked hard to be increasingly concise and logical, especially in research, which is the literature AI is heavily trained on, are being punished for writing that way. I wrote this entire paper aware of those tells, which is its own small proof of the point: knowing where AI gives itself away is exactly the kind of judgment the technology cannot supply for you.

The honest picture of what my own system cannot do: it still has edge cases. Refining a resume requires real trial and error. It does not make judgment-heavy steps much faster than from scratch. What it does is make the work more specific, more focused, and better than what I could have produced alone. The amplification is real. The replacement is not.

V. The Framework in Practice: What Companies Are Getting Wrong

Consulting firms, product companies, agencies, every organization building AI-powered services and delivering them to clients is facing the same three design questions that the Automate Amplify Preserve framework answers: what do we fully automate, what do we keep the human in the loop for, and what do we never let the system touch. The framework I built for a personal job search maps directly onto this. The parallel is not loose. The companies that get this right are building systems where the client's voice, values, and specific context are loaded into the system. Systems that have real human trust points at the beginning and end, where human support and explanation is easily accessible and backs the system that is more independent than the systems of the past, and more efficient. The companies that get it wrong are building systems that produce generic outputs at scale and calling it personalization. The uncanny valley problem that the 2016 SapientNitro paper described, replicating human service in ways that trigger discomfort rather than trust, is now the defining challenge in enterprise AI deployment. Think about an answering machine, now think about that being the only way to consume, or to have clients interact, in every step of those processes. That sounds like a future I want to avoid.

Prophet, for example, is a consulting and brand strategy firm that built MAIA, an internal AI system designed to power their own consulting work and support client co-creation. The parallel to what I built is deliberate. The presentation I brought to Prophet's AI Foundry team was titled Automate the Intake, Preserve the Thinking precisely because the philosophy that made my personal system work is the same philosophy that makes a client-facing AI system work. Both systems fail the same way when they are built wrong: they produce structurally correct, completely generic outputs that do not reflect the actual client, and the client can feel it immediately. And both systems succeed for the same reason: when the context architecture is specific enough, based on the company's proprietary methods and values, the outputs are genuinely useful rather than just plausible-sounding.

IBM Consulting announced at Think 2026 that it is now offering clients an "asset-based consulting service" to help them build and operate their own hybrid-AI platforms, with a stated goal of scaling AI "within their business context." (IBM, May 2026) That phrase is the context architecture problem stated in corporate language. Organizations that have tried to deploy generic AI tools and gotten generic results are now discovering that the hard part is building a system that actually knows their specific business.

The most advanced area of AI consulting in 2026 now involves building enterprise agents: systems that can reason, retrieve knowledge, use approved tools, and complete multi-step tasks. The failure mode at scale is identical to my failed first version. When those systems are not loaded with the specific context of the client they are serving, they produce statistically plausible outputs that are wrong for the specific situation. A customer service agent that does not know the company's actual values and policies will handle escalations in ways that feel subtly off. A research agent that does not have access to the firm's specific data and framing will produce analysis that could have come from anywhere. The token tumbler produces the statistical average. What the client needed was something specific to them, and something that they can work with, backed by someone who is checking everything along the way. The goal is to avoid a system with an input,

tons of steps in the dark, and the acceptance of an output that is simply just taken as the truth.

The Automate Amplify Preserve framework applies to every layer of this. At the workflow level, companies need to decide what they fully automate, what they amplify with a human in the loop, and what they preserve entirely for human judgment. Client strategy sessions should not be fully automated even if AI can produce a plausible-looking strategy document. Brand voice guidance cannot be generated without deep knowledge of who the brand actually is. Executive communication requires human leading, not technology leading, even when AI is structuring the final output. And the moments of genuine client relationship, the conversations where trust is built or lost, where a client's specific situation is heard and responded to personally, those are Preserve. Automation at that layer does not just produce a worse product. It damages the relationship in ways that are difficult to recover from.

There is a straightforward test for whether any AI deployment is in the right tier: put yourself in the shoes of someone receiving the output, who receives thousands of outputs like it every day. Does this feel like it was made for you specifically, or does it feel like it was made for the statistical average of someone in your situation? If the latter, the deployment is Automating something that should have been Amplified, or Amplifying something that should have been Preserved. The whole point of the framework is that the correct answer is different for every step of the workflow, and that getting it wrong in either direction has real costs.

I built a system that navigates these tradeoffs for one person across one process. The companies doing this work at scale are navigating the exact same tradeoffs across dozens of clients and hundreds of workflows. The skill that matters is the same: knowing which steps require what level of human involvement, building context architectures specific enough that the outputs reflect the actual client rather than the generic average, and designing human review checkpoints into every step where the cost of a wrong output is high. That is what I want to build professionally.

VI. The Broader Ripple: The Bubble Debate, Physical Limits, and What Everyone Is Missing

The AI bubble debate has been running for three years and the positions have not really moved. One camp says we are in a speculative frenzy indistinguishable from dotcom 2000, that valuations are untethered from revenue reality, and that the correction is coming. The other camp says this is fundamentally different, that the demand is real, the revenues are real, and the infrastructure buildout is being pulled forward by actual usage rather than speculation. Both camps have a point, and honestly both are a little too sure of themselves. Neither is fully right.

The dotcom comparison is instructive but imperfect. In 2000, telecom companies built massive fiber-optic infrastructure in anticipation of demand that was not yet there. Utilization rates sat below 3%. Revenue lagged investment by years. Infrastructure was funded through debt. The current AI cycle looks different in

several important ways: OpenAI and Anthropic are generating over \$20 billion in combined annual recurring revenue. More than 90% of new data center capacity is already committed before construction starts. Nvidia reported \$57 billion in revenue in a single quarter. The key beneficiaries are cash-flow machines, not speculative startups. (Redpoint Ventures 2026 Market Update; INSEAD Knowledge, 2026)

But there are also real risks the bull camp tends to dismiss. One Capital Economics economist concluded in March 2026 that the AI stock bubble has already partially deflated. Some of the circular financing structures are uncomfortable: Nvidia invests in OpenAI, OpenAI uses that capital to buy Nvidia chips, and Nvidia is guaranteed to be the supplier for that expansion. INSEAD researchers noted the distinction between "flywheel" and "house of cards" comes down entirely to whether real end-user demand is being generated, or whether money is moving in circles. (INSEAD, 2026) Experts at Fortune's 2026 Crystal Ball were unambiguous: "2026 marks the end of AI novelty purchasing. Companies will only pay for AI that increases revenue, reduces churn, or automates real work." (Stevie Case, Vanta, Fortune 2026)

My honest read on this, for whatever it is worth: the technology itself is not going anywhere. The internet did not disappear after the dotcom crash. It stabilized, matured, and eventually transformed everything. Amazon and Google survived and became the defining companies of their era. Something similar is likely here, and I really do think at least some level of this is here to stay. The hype around specific companies and valuations will correct. The underlying capability, the ability to process language, generate code, reason through problems, retrieve and synthesize information at scale, is genuinely useful in ways that do not disappear when the bubble corrects. What changes is that the bar moves from impressive demos to measurable ROI, which is exactly where it should be.

One thing both camps consistently underweight, and it connects back to the incentives problem from the start of this paper: the same companies warning about the dangers of AI are also raising capital at valuations that require the technology to be transformative. The urgency is not just narrative, it is structural to how these companies are financed. A bubble needs a story loud enough to justify the capital, and both the fear and the hype keep that story loud. The most useful thing you can do with any claim about AI, including a claim about whether it is a bubble, is ask who benefits from you believing it.

There is also a layer of this that the bubble discourse almost entirely misses: the physical world the models depend on. Everyone is focused on making the models better. Bigger context windows, more parameters, faster reasoning. The part of the story that is not getting nearly enough attention is the physical infrastructure these models require, and specifically, how badly that physical world is struggling to keep up.

Kevin O'Leary put it plainly at Data Center World 2025: "We're not just building data centers, we're building energy platforms around AI." His proposed Wonder Valley project is a 7,000-acre AI compute campus in Alberta powered by stranded natural gas, designed as a vertically integrated off-grid energy platform precisely because connecting to the public grid is no longer feasible for projects of this scale. He estimates returns of 14 to 20 percent on data center real estate while warning that 50 percent or more of announced US

data centers will not be built. Not because the demand is not there. Because there is no power to run them. (Data Center Frontier, 2025)

The constraint is well documented. The Uptime Institute identifies power as the single defining constraint on data center growth globally in 2026. Nearly 50% of all global data center projects scheduled for completion this year face delays directly attributable to power supply limits and grid shortages. Morgan Stanley forecasts a 49 GW power shortfall in the US alone through 2028. And the bottleneck is not just generation capacity. It is transformers. A five-year backlog on high-voltage grid transformers is effectively killing half of America's planned 2026 AI data centers. Pre-2020, lead time on a high-power transformer was 24 to 30 months. Today it stretches to five years. Electrical equipment is under 10% of total data center cost and 100% of the constraint. (Sightline Climate / Tech Fund, 2026)

"The defining risk for AI data center expansion has shifted from computational efficiency to the physical availability of grid-scale power, creating a primary bottleneck for commercial growth." EnkiAI / AI Data Center Grid Strain, April 2026

Data centers can typically be built in one to two years. Getting a grid connection takes two to ten years in the EU, and multiyear backlogs now exist across Northern Virginia, Texas, and other concentrated markets. (Enlit World, 2026) The result is that AI rack density is running ahead of everything else. GPU racks supporting AI training workloads exceeded 40 kW today, trending toward 85 kW, with 200 to 250 kW per rack projected by 2030. The electrical infrastructure to support that density takes a decade to build. That mismatch is the story most people covering AI are not telling.

The price impact is already visible. PJM capacity market clearing prices for 2026-2027 reached \$329 per MW, over ten times higher than \$28 per MW two years prior. Data center growth is explicitly identified as a major contributing factor. Power purchase agreements have become as strategically important as GPU procurement. Microsoft signed a 150 MW wind power deal with Iberdrola in Spain specifically to secure dedicated power for AI infrastructure. Companies are now building "bring your own power" (BYOP) campuses with on-site natural gas generation because waiting for grid connection is not commercially viable. Meta's campus near Columbus, Ohio went this route in 2025. It is worth noting that this means companies marketed as clean and forward-looking are increasingly powered by on-site fossil fuel generation.




Figure 4. Global data center electricity demand: actual and projected. The trajectory is steep in any scenario. The physical infrastructure to support it is not being built at the same pace. Sources: IEA Energy and AI Report (April 2026).

Water is where the criticism becomes most substantive. Some of the people closest to me object to AI primarily on this basis, and on this point I think they are closer to right than most people in the tech space want to admit, because again, that is in their self interest. Data centers use water to cool their hardware. Despite tons of varying findings, generating a 100-word AI response seems to evaporate roughly one bottle of water from cooling towers. (UC Riverside) Large data centers can consume up to 5 million gallons per day. US data centers directly consumed 17.4 billion gallons in 2023, projected to reach 38 to 73 billion gallons annually by 2028. (MOST Policy Initiative, 2026) Cities near data centers have seen temperatures rise 1.3 to 4 degrees Fahrenheit within a third of a mile of the facilities, from thermal plumes off air-cooled condenser arrays. (Arizona State University, 2026) A proposed data center near Applecross in Downingtown, where I live, has prompted local opposition and conversation about a planned protest at a city meeting. That response is not unreasonable. Residents are responding to documented effects in places like Phoenix and seeking to avoid the same outcome. I am an advocate for new technologies more than most, but I do believe that over generations we do not really think about the repercussions and tradeoffs of these advances as much as we should until it is extremely urgent to do so.

The accurate framing on water is that per individual interaction the use is small, but scaled to billions of daily queries it becomes a serious resource constraint, particularly in the arid climates where these facilities increasingly cluster. The opposition is not wrong about the underlying concern, and it is too often dismissed by those who treat the economics as justification enough. Up to 85% of the water data centers consume evaporates entirely rather than returning to the local supply. That is not a recoverable loss.

The O'Leary investment thesis is actually a useful lens here even beyond real estate returns. His point is that while everyone chases the obvious market: the model companies, the GPU manufacturers, the software platforms, the infrastructure plays are structurally more durable because they are physical constraints that cannot be solved with a software update. Power grid infrastructure, utilities, transmission capacity, cooling

technology, water recycling systems. Not glamorous. Structurally necessary. And unlike the model companies competing against each other on capability benchmarks, the infrastructure bottlenecks are supply constrained by physics and regulatory timelines, not by talent or compute.

The labor market reflects the same disorder. In several cases, the workers being cut are the wrong ones. Software engineers, who are more relevant in an AI world, are being cut while other more replaceable departments are sometimes kept. This is simply because the discourse is around the obsolescence of coding, however these same people forget to recall that coding is the foundation of the entire AI system to begin with. People who truly know how to code are better at prompting, better at context engineering, better at connecting and working with AI. Even if the systems are replacing their busy work, they can now focus much more on polishing those systems, or on system design, and focus on the intersections that expand their technical expertise into client-facing work or solutions. With the technical people gone, AI software updates are causing entire systems to fail and companies are trying to rehire the people they let go. Forrester predicted half of AI-attributed layoffs would be quietly reversed. The token pricing crisis adds more pressure. Anthropic's shift from flat-rate to per-token billing is removing budget predictability from enterprise teams at exactly the moment those budgets are under pressure from infrastructure costs. Some companies are now building systems to monitor and throttle AI consumption the same way they monitor cloud spend. The efficiency gains are real. The costs, physical and financial, are also real.

VII. What AI Cannot Touch: Relationships, Voice, and the Disconnection Crisis

Preserve draws a harder line than the rest of the framework, and it reaches past workflow design. It is not just a box in a diagram. It is a stance on which parts of work and life should stay human, no matter what the technology can do.

There is a broader pattern here, and it is one I notice more the older I get. People are increasingly absorbed in work, and outside of work they anchor their sense of self-worth to online comparison, measuring themselves against curated versions of people who, day to day, are nothing like what they present. Most of us now spend more time with devices than with family or the people closest to us. The research on the consequences is not ambiguous. A 2025 APA survey found that roughly half of all adults report feeling isolated, left out, or lacking companionship. Academic research places moderate-to-severe loneliness at around 37% of Americans. The US Surgeon General has declared loneliness a national epidemic. The technology industry's response has been to offer AI companions.




Figure 5. Loneliness, AI companionship, and the disconnection paradox. Sources: APA (2025); arXiv chatbot study, Fang et al. (2025); faspsych.com (2025).

Heavy daily use of AI chatbots, similar to extensive computer use, correlates with greater loneliness, more dependence, and reduced real-world socializing. A four-week randomized controlled trial found that only participants who texted with a real human peer showed a significant reduction in loneliness. Those interacting with the chatbot showed levels statistically indistinguishable from those who simply journaled. AI companions promise connection and deliver engagement. Those are not the same thing.

There is a whole sector of people in the tech space who are extremely motivated and inspired by the extra productivity, or the feeling of volume, from AI to a fault. They are addicted to automating their lives, and instead of stopping when it makes their lives easier, they find themselves working more than they ever have before simply because they are able to, and less critical thinking is required over extensive time periods to get the same amount of work done, so they can do it for longer.

Mrinank Sharma, head of Anthropic's Safeguards Research team, resigned in February 2026 with a letter warning "the world is in peril" from "a whole series of interconnected crises." His final project at Anthropic was specifically on understanding how AI assistants could make us less human or distort our humanity. (CNN, 2026) A researcher departing OpenAI the same week warned the technology has "a potential for manipulating users in ways we don't have the tools to understand, let alone prevent." These are people who built these systems, from the inside, asking whether building them was the right call. The thing that stays with me is how often the people closest to the frontier, the ones who would know, end up saying some version of the same quiet thing: spend time with the people you love while you still can. That is not a marketing line. It is coming from the people with the least incentive to say it, and that is exactly why it lands.

"We appear to be approaching a threshold where our wisdom must grow in equal measure to our capacity to affect the world, lest we face the consequences." Mrinank Sharma, Head of Safeguards Research, Anthropic | February 2026

The things worth prioritizing in this period are the ones that are neither scalable nor automatable. Relationships matter first, because a worker who is purely technical and fully remote is more easily replaced than one who has built genuine trust with the people around them. Physical health matters because the demand for cognitive output or staying power is higher than it has ever been, and the body still determines how sustainable that output is. Creative work pursued for its own sake rather than for productivity has value precisely because it resists optimization, and stands out in genericism. The same logic applies to direct human expertise. In most cases, a single conversation with a knowledgeable person yields more in minutes than extended trial and error with an automated system, and has a feeling of assuredness that is unreplicable elsewhere. Those interactions are steadily being replaced with systems that are cheaper but less effective, and the tradeoff is rarely made deliberately.

My real concern, the one underneath all of this, is complete autonomy. Letting AI systems roam and figure things out on their own, and approving whatever they decide blindly, is a different category of risk than a bad output. If we keep handing more and more authority to systems whose reasoning we cannot fully audit, then in time they effectively run things. For dramatization sake, let's say one direction: they stay rational and we get something like the WALL-E version of the future, leading lives based around instant gluttonous gratification, void of purpose, while the machines delegate all the structure. Or another: what if they reach a point where they decide humans are unnecessary, and we handed them the positions and the authority to act on it. I am not saying that is imminent, or realistic. Simulations of AI worlds have seen a vast range of outcomes. I am saying the design choice that prevents it is not exotic: there should always be a human approving the steps that matter. Leaving consequential decisions to systems we cannot fully audit is the path toward outcomes that range from unfortunate to irreversible. The human in the loop is not there because AI is always wrong. It is there because the decision that matters has to belong to someone who can be held accountable for it.

VIII. A Personal Philosophy for the AI Era

Four principles. All arrived at through building something specific, watching it fail, and diagnosing what went wrong. Not abstract. They came from work.

"The quality of AI output never exceeds the quality of the context it is working from."

Generic inputs produce generic outputs. This is an architectural insight, not a prompting insight. Building the context architecture, what the system knows, how it is structured, how it is maintained and updated, is the hard part. The tools are easy. This applies to personal workflows and to enterprise deployments equally. A consulting firm's AI system that does not know the client in depth will produce outputs that sound like consulting. It will not produce outputs that sound like that client's specific situation.

"Specificity is the entire skill."

Ask the AI to ask you questions before it produces anything. Pressure test every output. Push back when something feels generic or off. Be deliberately redundant on rules. Check everything. Rewrite everything that does not sound like you. The people who will use AI best are not the ones who accept outputs fastest. They are the ones who have developed enough taste to know when something is right.

"Where AI belongs and where it does not is a design decision, not an afterthought."

Automating formatting is safe. Automating relationship building is not. Automating intake is safe. Automating judgment calls about fit and voice is not. Every workflow decision, whether you are designing a personal system or deploying a client-facing one, reflects a judgment about what AI can be trusted to do and what it cannot. That judgment, applied systematically, is the system. As a society, we need to make these decisions consciously, at scale, now.

"The goal is not to do less work. It is to make sure the work is genuinely yours."

AI gives everyone more capacity to affect the world than at any prior point in history. The question of whether that capacity is used well is entirely a question of wisdom, judgment, and values. None of which the technology provides. The people who shape this well are the ones who understand where AI helps and where it breaks, who preserve the human moments that actually matter, and who design their work with the specific goal of amplifying what is distinctly human rather than replacing it with something that only resembles it.

I want to be genuinely optimistic about this, and I mostly am. There is a version of the AI era where the tools amplify what is specific and human, where the time freed up by automation gets directed toward the relationships and creativity that actually matter, where companies that deploy AI to their clients do it in a way that makes those clients feel more seen rather than less. Getting there takes deliberate design choices at the level of the individual deciding how to use a tool, then at the level of the company deciding what to ship, and then at the level of policy deciding what to allow. None of those layers fixes the others. The default, if nobody makes those choices on purpose, is generic output at scale and a slow erosion of the things that made the work worth doing in the first place.

I am not certain which version we get. That is the honest position, and anyone telling you they know the future is selling you something. What I am certain of is narrower. The people worth paying attention to right now are not those sellers or the ones predicting the most dramatic collapse. They are the ones building specific things, testing them against reality, saying plainly when something does not work, and putting a human in every step where being wrong actually costs someone something. I would rather be one of those people than be right about the forecast.

So here is where I will leave it. The race right now is to see who can benefit the fastest, who can automate the most, who can claim the most ground before everyone else. I think that is the wrong race. The harder and more important thing is to stop, look at what we are actually building, and figure out how to deploy it in the right way before the choice gets made for us. AI gives us leverage. It does not give us direction. Direction is still ours, and it is the one thing we cannot afford to automate away. The consequences of getting this wrong are severe enough that we have to be proactive instead of reactive, and the cost of waiting until it is obviously urgent is a cost we do not get to take back.

None of this comes from being against AI. I am the opposite of that. I embrace it without hesitation, I use it every single day, and I think it is here to stay and capable of fostering developments we cannot yet imagine. That is exactly why I care about this. The one thing I want people to hold onto, through all of it, is their own ability to think, and to think creatively. Stay smart. Dive into the weeds of the things you are actually passionate about. Use your own creativity instead of outsourcing the whole of your thought to a machine. And build the systems that protect that human element rather than erase it. I want the future where this technology amplifies us. I am just not willing to sleepwalk into the one where it quietly replaces the part of us that made the work worth doing in the first place.

"By far, the greatest danger of Artificial Intelligence is that people conclude too early that they understand it." Eliezer Yudkowsky

"Artificial intelligence is not a substitute for human intelligence. It is a tool to amplify human creativity and ingenuity."

"AI gives us leverage. Ethics gives us direction."

"The consequences of AI going wrong are severe, so we have to be proactive rather than reactive."

References

Context Engineering: A Practitioner Methodology for Structured Human-AI Collaboration. arXiv:2604.04258, April 2026.

Murray, R. & Kopen, D. Experience Automation: Using Technology to Amplify the Human Touch. SapientNitro Insights, 2016.

Karpathy, A. Introduction to Large Language Models. YouTube lecture. andrej.com, 2024.

Russell et al. StoryScope: Structural Fingerprints of Human and AI Narrative. arXiv:2604.03136, 2026.

IBM Consulting. IBM Expands AI Capabilities to Accelerate Enterprise Transformation. Think 2026, May 2026.

Metaintro. Will AI Replace Recruiters in 2026? metaintro.com, May 2026.

Graham, J. et al. Artificial Intelligence, Productivity, and the Workforce. Duke / Federal Reserve CFO Survey, March 2026.

International Energy Agency. Key Questions on Energy and AI. iea.org, April 2026.

Sharma, M. Resignation letter. Anthropic. February 9, 2026. Reported CNN Business, February 2026.

Forrester Research. Predictions 2026: The Future of Work. December 2025.

Oxford Economics / Fortune. AI Layoffs: Corporate Fiction or Structural Shift? January 2026.

O'Leary, K. Keynote: Wonder Valley and the Great AI Pivot. Data Center World, March 2025.

American Psychological Association. 2025 Stress in America Survey.

Fang et al. Randomized controlled trial of AI chatbot companions and loneliness. arXiv, 2025.

Shortlistd. 10 Hiring Predictions for 2026. shortlistd.io, December 2025.

SSRS / Edison Research. Half of Americans Using AI Chat on Weekly Basis. March 2026.

Pew Research Center. Key Findings About How Americans View Artificial Intelligence. March 2026.

OpenAI via TechCrunch. ChatGPT Surpasses 900 Million Weekly Active Users. February 2026.

DemandSage / Botpress. Global AI Chatbot User Statistics 2026. 2026.

Ofcom. Generative AI Use Among UK Online Adults. 2026.

Deloitte. Generative AI Adoption by Generation. 2026.

SurveyMonkey. AI Trends by Generation: 25,030 US Adults. February 2026.

Edubrain via Fast Company. Which Generation Uses AI the Most. February 2026.

Investing.com. The AI Token Pricing Crisis Behind OpenAI and Anthropic's Revenue Race. May 2026.

Let's Data Science. Anthropic and OpenAI Face AI Token Pricing Crisis. May 2026.

Redpoint Ventures. 2026 Market Update: AI Cycle Built on Real Demand. 2026.

INSEAD Knowledge. Are We in an AI Bubble? February 2026.

Fortune / Term Sheet. Crystal Ball: Will the AI Bubble Burst or Balloon in 2026? January 2026.

CityAM. Why the AI Boom Cannot Be Compared to the Dotcom Bubble. March 2026.

Capital Economics (Higgins, J.). AI Stock Bubble Has Already Partially Burst. March 2026.

Uptime Institute. Five Predictions for 2026: Power as the Defining Constraint. 2026.

Sightline Climate. Analysis of 2026 US Data Center Project Pipeline. February 2026.

Morgan Stanley. AI Infrastructure Power Forecast. 2026.

Enlit World. AI Growth vs. Power Availability: A Perfect Storm in 2026. January 2026.

EnkiAI. AI Data Center Grid Strain: Power Halts Growth in 2026. April 2026.

Tech Fund / TechInvestments.io. Power Bottlenecks and the AI Data Center. May 2026.

arXiv. Electricity Demand and Grid Impacts of AI Data Centers: Challenges and Prospects. 2026.

MOST Policy Initiative. Data Center Water Use. April 2026.

Arizona State University / Smart Cities Dive. Data Centers Raise Temperatures Up to 4 Degrees. May 2026.

University of California, Riverside. Water Consumption from AI Query Processing. 2024.

George Mason University College of Public Health. AI, Loneliness, and the Value of Human Connection. 2025.

SHRM. The State of AI in HR 2026 Report. April 2026.

InteliScience. Enterprise AI Consulting for Enterprises: How to Choose a Partner. April 2026.

This paper was written with assistance from Claude. The research, reasoning, opinions, and voice are the author's own.