

AUTOMATE THE INTAKE. PRESERVE THE THINKING.

A Portfolio Project by Graham Friedman

BS Statistics & Economics + MS Quantitative Economics · University of Pittsburgh

AI Workflow Systems · Claude · Claude Code · Anthropic API · Python

The Position and the Problem.

The Background

Statistics and Economics at Pitt, now enrolled in a Master's in Quantitative Economics. Deep training in how to frame analytical findings for audiences making real decisions, not just reading models. Data science coursework, coding experience, research experience, and extensive hands-on time building with, auditing, and rebuilding AI systems.

The Gap in the Market

This is the first recruiting cycle meaningfully affected by AI at scale. Tools promising to automate applications produce outputs that could have come from anyone. Not aligned to who you are, fabricated at scale, optimizing for volume not fit, leaving no room for building relationships in the fields you really want to be in. They are not just unhelpful. They are a major contributor to the noise recruiting teams are still trying to manage, and the main contributor to the noise they promise to help you fight through.

The Core View

AI should handle the mechanical, repetitive, non-relational work. Analysis still requires human judgment to evaluate whether a finding holds and what it means for the specific situation. Anything involving a client's trust, their relationship with your firm, their actual needs being heard. That stays fully human. AI amplifies the thinking. It does not replace it. As companies build systems to put directly in front of customers, they need to be careful to choose which steps of their process customers would want to have preserved human judgement and connection. Uniqueness will stand out and be sought out. Creativity is more important now than ever before.

The Response

Build a system that finds the balance between efficiency and the human elements that actually matter. Speed on the mechanical work. Preserved judgment on everything requiring your actual voice, your real story, and your genuine fit with the people you are reaching out to. In a job obtaining setting, showing why your framing specifically aligns with this type of work, and advertising yourself in that unique way, will still being efficient in that process.

Where I Stand on Operating In and On AI Systems.

Not an engineer. Not a passive user. Something more useful than either in a consulting context.

I build in the interface, not the infrastructure.

My work with AI happens at the application layer: designing prompts, structuring context, building workflows, and deciding what to automate versus preserve. I do not train models or architect infrastructure. I figure out how to make existing systems actually work for specific problems with specific constraints.

I know where AI breaks before it causes damage.

A year auditing AI statistical outputs at a data annotation company taught me exactly how models fail: systematic misclassification, spurious pattern recognition, confident wrong answers that look right on the surface. That experience is worth more in a consulting context than knowing how to build a model from scratch. The risk is not that AI fails obviously. The risk is that it fails in ways that are hard to see.

I treat AI as a thinking partner, not an answer machine.

Every workflow I have built assumes the AI needs to be questioned. I ask it to ask me questions before it produces anything. I pressure test outputs. I push back when something feels generic. I read everything out loud before it goes anywhere. The system does not replace judgment. It makes judgment faster and better informed.

I think about deployment, not just capability.

Building a capable AI system is the easy part. Figuring out how to deploy it in a way that works for real people with real constraints, where they still feel like their actual needs are seen and taken care of, is where the hard problems are. That is the intersection I want to work at. Not what the system can do in isolation. What it actually produces for the person on the other end.

The Foundation. What the System Actually Knows.

Three layers: project files permanently loaded in Claude Project, master prompts pasted per session, and the Notes app underneath everything.

PROJECT FILES: PERMANENTLY LOADED IN EVERY SESSION

Master Background

9-part complete professional identity document

Full Pipeline & Trackers

Interaction history with every contacted company. Depth beyond a full CRM.

Writing Rulebook

Voice governance built from personal writing samples

My AI Experience

What I've built, how I use AI day-to-day, where I've seen it fail, and opinions on the future

Full GitHub Description

Complete profile and all project READMEs in one document

Full LinkedIn Description

Complete profile with cross-platform discrepancies flagged

Base Resume PDF

One-page base resume. System reads visually before every tailoring step.

MASTER PROMPTS: PASTED IN WHEN THE STEP NEEDS THEM

How to Use Claude

Extensive living reference on how to actually work with Claude. Covers methodology, prompting principles, and workflow design. Anchored around Anthropic's five official courses.

Pipeline Management

Governs full pipeline reviews. Three pipeline categories, exact gameplan format, prioritization order, dead thread rules, reactivation flags.

LinkedIn Outreach

Governs all networking drafts. Specific real reference required, direct ask matched to seniority, under 200 words for InMail or 300 characters for connection note.

Resume Tailoring Machine

Go through all coursework, skills, competencies, project details. Output only changes with before and after text. Never increase length.

NOTES APP: THE MESSY WORKSPACE BEHIND EVERYTHING

Every project file lives there as an exact backup. But Notes is also where everything else lives: stream of consciousness, raw thoughts, drafts in progress, ideas that have not become polished yet. My whole life is basically on my Notes app. Which means when something changes in how I think about a role, a company, or my own positioning, it shows up in Notes first before it ever becomes a polished file. The system stays current because Notes is current. Any time a project file changes, Notes updates first, the full file gets copied, and it gets repasted into Claude.

These 7 files also load as the system prompt for the JD Analyzer, which runs separately through Claude Code. See next slide

The JD Analyzer. The Only Fully Automated Step.

Lives in Claude Code, a terminal environment, not a chat interface. A Python script calling the Anthropic API directly.

HOW IT WORKS

INPUT

Paste the raw job description into Claude Code terminal.

CONTEXT

Full 7-file background loads as system prompt via direct Anthropic API call. All 7 files loaded simultaneously as a single system prompt.

ANALYSIS

Script returns a structured 5-section brief: top skills emphasized, core themes, best projects to lead with and why, recommended framing angle, red flags or gaps to address.

OUTPUT

Complete strategic brief produced in under 60 seconds. The reasoning is tied to the specific JD language, not general impressions. Raw output exactly as it appears in the terminal. No editing, no cleanup.

NEXT

Human reviews brief, makes all judgment calls, proceeds to application. System informs, never decides.

WHY THIS IS THE ONLY AUTOMATED STEP

Specific to the person

The system prompt IS the person in full, which means generic framing is structurally impossible.

Replaces intake, not judgment

One script call replaces 20 minutes of careful manual JD reading. The script does not replace that judgment. It does the intake so human judgment is reserved for what requires it.

Context boundary is deliberate

Claude Code is the wrong tool for anything requiring genuine context: conversation history, pipeline status, relationship nuance. The Claude Project holds all of that. Claude Code handles what is mechanical. The Claude Project handles everything requiring knowledge of the person. That boundary was a design decision, not a limitation I worked around.

Best projects with reasoning




The brief tells you which projects to lead with and exactly why, tied directly to the JD language. Not based on what sounds impressive generally. Based on what this role is actually looking for.

Cost is real, but not the main reason

API calls have a real cost per use. That is part of why this is the only Claude Code step. But the deeper reason is the context boundary above. A script cannot hold relationship history, pipeline nuance, or prior conversation context. The Claude Project can. That is not a cost tradeoff. It is a fundamental architectural difference.

The brief becomes the strategic input for every step that follows: resume tailoring, portal answers, outreach framing, and in interview prep.

The System in Action.



A walkthrough of the system as it actually exists. Notes app open to the How to Use Claude file, then switching to Claude Code terminal running the JD Analyzer, then opening the Claude Project file list and scrolling through the Master Background file to show the depth of context loaded in every session.

JD Analyzer terminal output. Raw structured brief returned in under 60 seconds from a single script call. No editing. No cleanup. Exactly as it appears.

Second terminal screenshot showing the full 5-section brief structure: top skills emphasized, core themes, best projects to lead with and why, recommended framing angle, red flags or gaps to address.

A NOTE ON THE ADVICE FILE.

You will notice an additional file in the project called All Advice on Job Obtaining Processes. This file is not a core system file and is not required for the system to function. It is a living reference document compiled from transcribed videos, reels, and content covering career strategy, networking tactics, interview frameworks, and AI tools from a wide range of external perspectives. It gets added as supplementary context on steps involving strategic judgment calls: outreach method selection, cover letters, niche situation advice, pipeline management, and anywhere a broader external perspective adds value beyond what the personal background files contain. Think of it as a curated second opinion layer rather than a foundational system component.

The Background File. A Deep Dive Into Contents.

9 distinct parts. Primary identity document powering every resume, portal answer, cover letter, and interview answer. Generic inputs produce generic outputs. This file exists to make generic output structurally impossible.

Part 1: Identity

Name, contact, LinkedIn, GitHub, target roles across financial analytics, quantitative analytics, AI analyst, sports analytics, consulting analytics, due diligence. Target company categories. Positioning summary.

Part 2: Education

Two-layer structure: resume-ready formatting with all rules locked in, plus every course taken with exact grades so the system selects the 6-8 most relevant per role. MS always listed first. Subline wording locked verbatim.

Part 3: Experience

Data Annotation with 4 sub-sections: exact resume bullets, full context on how the role was obtained through multi-stage statistical assessment, complete 3-step diagnostic framework, full STAR example (retail dataset, 50 stores quarterly over 3 years, scene-by-scene account of catching panel data misclassification, multicollinearity, heteroskedasticity), 5 explicit transfer connections mapping the experience to QA, operational controls, audit, model risk, and compliance analytics. Air Quality with sensor readings, failure tree methodology, verification sequence, emergency service STAR story. Lawn Care documented as a single-line entry.

Part 4: Projects

Six analytics projects each with complete technical depth: exact R packages and Stata commands, dataset sizes, model parameter values, CV results, RMSE and R-squared values, statistical test statistics and p-values, full justification for every methodological decision. Used Car documents GAM chosen over Random Forest despite 0.006 RMSE disadvantage. Ravens Crime documents 80 hours of independent freshman work and counterintuitive playoff finding. NBA documents 3 heteroskedasticity tests. Lebanon documents 7 simultaneous macro frameworks and \$18B in fictitious seigniorage. Ravens Economic Activity documents honest null result with specific limitations. AI Workflow Systems contains exact verbatim resume bullets used without modification, since this project is the system itself and the framing must be precisely controlled.

Parts 5 and 6: Skills and Differentiators

Complete technical skills list ordered by JD relevance, max 12 on resume. Full data competencies bank. 10 specific differentiators in priority order with framing guidance for each, led by completing both a BS and MS in 4 years total, Magna Cum Laude, Certificate in Sports Analytics.

Parts 7 and 8: Personal Story and Honors

Origin story, broken legs story from 8th grade, Lamar Jackson connection, study abroad relationship context, honest weakness framing, specific end-of-interview questions for interviews only, never on a resume. Integration Bee Qualifier, Magna Cum Laude, Summa Cum Laude high school, Applecross Basketball Association founder, volunteering history, language proficiencies.

Parts 9 and 10: Resume Rules and Portal Answers

Every formatting rule: one page only, banned and preferred action verbs, two-second recruiter test, no em dashes anywhere in any output ever. Paste-ready portal descriptions in 4 versions: Operational Controls/QA, Data Analytics/BI, Project Management/Client Ops, AI/Consulting/Due Diligence. Full spoken-voice behavioral answers and pre-built one-sentence framings for specific interview contexts, all tweaked by role.

Built as a living document with explicit update rules. New roles, projects, and credentials slot directly into their sections without rebuilding anything. Updated in Notes first whenever anything changes, then repasted into Claude in full.

The Remaining Files. Summarized Contents.

1 — Full Pipeline & Trackers

Complete interaction history with every company contacted throughout the search. For every company: full chronological history of every action taken with exact dates, all contacts with titles and how they were found, every message sent verbatim, every response received. Strategic position assessment explaining where things stand and why. Phased action plans with specific triggers for each phase. Explicit behavioral rules: what to do and what never to do. Depth of a full CRM. Specifics kept private.

3 — My AI Experience

Deliberately unpolished first-person reasoning across 5 questions: what was built and how it works, ownership statement on who actually built the automation, how AI is used in research day-to-day, honest view on where multi-agent AI is heading including concerns about AI safety and renewable energy infrastructure as a durable investment, and specific moments where AI genuinely surprised including the first time Claude built PowerPoint slides from a description, the failure of Version 1, and the teleprompter deploy. Authenticity was the design goal. Not cleaned up after initial answers.

5 — Full LinkedIn Description

Complete current LinkedIn profile copied in for outreach tailoring. Headline, About section, all experience descriptions, education including all honor societies and performing arts organizations, skills section. Known cross-platform discrepancies flagged explicitly. For example, HVAC role title differs between LinkedIn and resume and both sites are documented with the reason.

2 — Writing Rulebook

Voice governance extracted from personal writing samples so Claude learns actual voice rather than approximates it. The gap between AI output and human voice is a design problem, not a prompting problem. This file is the solution. Section 1: sentence-level patterns: long sentences accumulating detail through parentheticals, burstiness between very long and very short, 'so' to start conclusions, 'also' to add things, specific hedges, no passive voice, thinking while writing not before. Section 2: comprehensive banned list: Moreover, Furthermore, Leveraging, Utilizing, Spearheading, Seamless, Passionate about, Proven track record, Excited to share, rule of three, participial phrase closers. Section 3: 8 sentence-level rules. Section 4: detailed rules for 6 output types including LinkedIn posts, networking messages, resume bullets, and interview answers. Section 5: 3 universal rules. Section 6: 12-question self-check before any output.

4 — Full GitHub Description

Complete GitHub profile and all project repository READMEs compiled into one context document. Ensures consistent positioning across every application without drift between what GitHub shows and what the resume claims. Profile bio. Repository overview emphasizing model selection, interpretability, and analytical judgment over raw predictive performance. Full README text for every project including exact research questions, dataset overview, complete methods, core findings with specific statistics, policy implications, and honest notes on reproducibility.

6 — Conversation Notes: 4 Master Prompts

How to Use Claude: extensive living reference on methodology, prompting principles, and workflow design anchored around Anthropic's five official courses — the entry point, not the full depth of the document. Pipeline Management: defines 3 pipeline categories, exact gameplan format with Full History, Current Status, Strategic Position, numbered Phases, Behavioral Rules, and Strategic Summary, plus prioritization rules, closing rules, and reactivation rules for the future full-time cycle. LinkedIn Outreach: requires specific real reference from the person's profile, connection to most relevant experience, direct ask matched to seniority, under 200 words for InMail or 300 characters for connection note, banned openers explicitly listed. Resume Tailoring: go through all coursework, skills, competencies, and project details, eliminate redundancy across bullets, explain project selection reasoning, never increase resume length, output only changes with before and after text.

How It Was Built. And How I Think With AI.

The depth of those files did not come from uploading a resume and asking Claude to summarize it. This process was applied to every file containing personal content. Not just the background file.

01 Upload Everything

Every piece of raw material provided upfront. Resume, project files, course history, work documentation. Nothing held back.

02 Ask Claude to Ask Me

Rather than telling Claude what to write, asked it to generate every question it needed before producing anything. Not clarifying questions about facts. A full interview. What did it feel like the moment the AI failure became clear. What was happening emotionally during the diagnosis. Why did that particular error feel satisfying to catch when nobody else had noticed it. What motivated every decision, not just what the decision was. What made each experience meaningful rather than just documentable. What would never appear on a resume but shaped how you think. Claude asked everything it could think of across every experience, every project, every file. Nothing was assumed. Nothing was skipped.

03 Answer in Unpolished Voice

Every question answered in full without cleaning up. Authentic reasoning was more valuable than a polished summary. The raw thinking is the point. Stream of consciousness, half-formed ideas, emotional texture, all of it went in exactly as it came out.

04 Preserve the Voice

Explicitly instructed to preserve actual language and structure while making output clean and specific. Goal was refinement, not replacement. Claude was not asked to rewrite. It was asked to make the raw answers clean and specific while keeping the structure and language intact. The difference between those two instructions is the difference between AI output that sounds like you and AI output that sounds like AI describing you.

05 Push Back

Flagged everything that did not sound right or was missing nuance. Asked Claude to revise. Repeated until every section was accurate and fully defensible. If a sentence felt even slightly generic it got challenged. The standard was: could this sentence have been written about someone else. If yes, it was not done yet.

06 Be Deliberately Redundant

Constantly reminded Claude of structural rules. Explicitly restated constraints in every prompt. Specificity is the entire skill. Generic inputs produce generic outputs. Every single time. The Writing Rulebook exists because repeating rules in every prompt is not enough. The rules needed a permanent home.

This was not a one-time build. Every significant conversation, new application, and new piece of feedback updates the files. The system improves because the context improves.

The Failure That Made the System.

VERSION 1: WHAT FAILED

- ✗ The first version was a reasonable architectural choice. Outputs were structurally correct. The failure was not immediately obvious.
- ✗ Built a Claude artifact calling the Anthropic API directly across 7 automated steps.
- ✗ Resume tailoring, portal generation, networking outreach, pipeline tracking, all automated sequentially.
- ✗ Outputs were structurally correct and completely generic.
- ✗ No contextual memory across sessions. No knowledge of what happened in prior conversations.
- ✗ Did not sound like the person. Sounded like a well-organized AI description of the person.
- ✗ Could not calibrate to tone, relationship history, or specific recruiter context.
- ✗ Scrapped entirely.



VERSION 2: THE REBUILD

- ✓ Failure was architectural, not technical. The system did not know the person.
- ✓ Rebuilt around Claude Project architecture holding all 7 files simultaneously across every session.
- ✓ Context architecture is the hard part. Tools are the easy part.
- ✓ Writing Rulebook extracted from personal writing samples so Claude learns voice rather than approximates it.
- ✓ Every output reviewed, read aloud, and rewritten where it does not sound right before anything goes out.
- ✓ Full conversation history maintained across sessions: system knows every prior interaction.
- ✓ The rebuild required starting over on the context architecture entirely, which is where the 6-step build process on the prior slide came from.
- ✓ **Generic output is now structurally impossible because the system prompt IS the person.**

The Complete Workflow. Preliminary Work.



AUTOMATE



AMPLIFY



PRESERVE

HOW THIS WORKS

The 7 project files are permanently loaded in the Claude Project and available in every session automatically. Claude sees all of them simultaneously without you doing anything. You still tell Claude explicitly which files are most relevant for each step because naming them activates focused attention on them. Conversation notes are master prompts pasted manually at the start of the session when that step requires them. They are not project files. Each company gets its own focused chat session within the broader Claude Project so conversation history stays coherent and specific to that company.

MORE ON HOW THIS WORKS

You have to tell Claude what to reference every time. For each step paste all the file titles it needs for that specific task. The project files load automatically but Claude still benefits from being told which ones are most relevant. Conversation uploads must always be explicitly pasted. Repastes like a job description are not always necessary if pasted recently within a single session, but redundancy is otherwise very important. The Writing Rulebook exists so less rewriting is required but pure AI voice sounds nothing like a real person. Rewriting is always still necessary regardless. Notes app is the source of truth. Any time a file changes, Notes updates first, the full file gets copied, and it gets repasted into Claude.

PRELIMINARY WORK

Everything built and maintained before the job search begins. These files form the foundation the entire system draws from.

1.1 Re-evaluate or update career trajectory (Amplify — Human Leading)

FILES: Background · Pipeline · AI Experience · Advice

PASTE: Updated career goals or constraints in plain text

CLAUDE IS TOLD TO: Reassess direction given updated goals. Identify which specific sections of the Background and Pipeline files need updating and rewrite those sections incorporating the new direction.

OUTPUT: Proofread output. Go into Notes, find those exact sections, replace them verbatim with Claude's output. Reupload the full files to Claude next session.

1.2 Update master background (Amplify — Human Leading)

FILES: Background · Writing Rulebook · AI Experience · Pipeline · Advice

PASTE: The specific section of the Background file being updated pasted directly. Reason for the update in plain text.

CLAUDE IS TOLD TO: Rewrite that specific section incorporating the new information following all Writing Rulebook rules. Do not touch any other section. Output only the rewritten section.

OUTPUT: Proofread output. Go into Notes, find that section, replace it exactly. Reupload the full Background file next session.

1.3 Update base resume (Amplify — Technology Leading)

FILES: Background · Writing Rulebook · Pipeline · Advice

PASTE: The specific resume section being updated pasted directly. Current base resume PDF uploaded to conversation. Reason for the update in plain text.

CLAUDE IS TOLD TO: Rewrite that specific section following every resume rule in the Writing Rulebook. Output only the changed section with before and after text. Do not output the full resume.

OUTPUT: Read output out loud. Rewrite where it does not sound like you. Go into docs, update the base resume, create updated PDF, upload to Claude Project.

1.4 Update GitHub (Automate)

FILES: GitHub · LinkedIn · Background · Advice

PASTE: Current GitHub section needing update pasted directly. Reason for change in plain text.

CLAUDE IS TOLD TO: Rewrite that specific GitHub section so it is consistent with the Background file and aligned with the LinkedIn profile. Output only the rewritten section.

OUTPUT: Proofread output. Make the change live on GitHub. Go into Notes, find the GitHub file section, replace it exactly. Reupload the full GitHub file next session.

1.5 Update LinkedIn (Amplify — Technology Leading)

FILES: LinkedIn · GitHub · Background · Writing Rulebook · AI Experience · Advice

PASTE: Current LinkedIn section needing update pasted directly. Reason for change with as many details as possible.

CLAUDE IS TOLD TO: Rewrite that specific LinkedIn section following the Writing Rulebook. Ensure consistency with GitHub and the Background file. Output only the rewritten section.

OUTPUT: Proofread output. Make the change live on LinkedIn. Go into Notes, find the LinkedIn file section, replace it exactly. Reupload the full LinkedIn file next session.

The Complete Workflow. The Application Process. Part 1.



AUTOMATE



AMPLIFY



PRESERVE

2. THE APPLICATION PROCESS

The full chronological workflow for every individual application. From finding the role through the cover letter.

A NOTE ON THIS SECTION

The JD Analyzer in step 2.2 is the only step that runs through Claude Code rather than the Claude Project chat interface. Every other step in this section runs through the Claude Project which holds full conversation history and context across sessions. The JD Analyzer brief produced in 2.2 becomes the strategic input for every step that follows. Jobalytics referenced in 2.3 is an external platform that generates a keyword match percentage and a list of missing keywords from the JD. Used alongside the resume tailoring step to ensure keyword coverage without forcing unnatural language.

2.1 Locate jobs (Automate)

FILES: Background · Pipeline

PASTE: Current search criteria and role types in plain text. Ask Claude to search for companies matching the target profile.

CLAUDE IS TOLD TO: Search for companies matching the target profile based on role types, industries, and location preferences in the Background and Pipeline files. Return a list of specific companies with open roles.

OUTPUT: Review the list. Check company job boards directly once targets are identified. Nothing needs to be saved back into the system.

2.2 JD fit assessment via JD Analyzer (Automate)

FILES: All 7 project files loaded simultaneously as system prompt via Claude Code

PASTE: Raw JD pasted directly into Claude Code terminal. Run JD Analyzer Python script. Script calls Anthropic API with full background as system prompt.

CLAUDE IS TOLD TO: Return a 5-section brief: top skills emphasized by the JD, core themes and role framing, best projects to lead with and why tied to JD language, recommended framing angle, red flags or gaps to address. Raw output, no cleanup.

OUTPUT: Review brief. Use as strategic input for every step that follows. Nothing needs to be saved back into the system.

2.3 Optimize and customize resume for JD (Amplify — Tech Leading)

FILES: Background · Writing Rulebook · Pipeline · LinkedIn · GitHub · Advice

PASTE: Resume Tailoring Machine master prompt. JD pasted directly. Jobalytics keyword score and missing keywords pasted directly. Current base resume PDF uploaded. Ask Claude to output full updated section for any bullet being changed.

CLAUDE IS TOLD TO: Go through all coursework, skills, competencies, and project details. Rewrite every bullet that can hit JD keywords naturally.

Eliminate redundancy within the same role. Recommend projects to include or drop with reasoning. Reorder coursework for JD relevance. Update skills section to lead with most relevant tools. Output only changes with before and after text. Do not increase length. Do not output the full resume.

OUTPUT: Proofread each changed section. Read each bullet out loud. Rewrite where it doesn't sound like you. Copy all changes into a new doc copied from the base resume and save for upload.

2.4 Fill out portal answers (Amplify — Tech Leading)

FILES: Background · Writing Rulebook · LinkedIn

PASTE: JD pasted directly. Specific portal fields that need answering. Direct Claude to write in paragraph longer form where it makes sense.

CLAUDE IS TOLD TO: Write answers to each portal field drawing from the full background. Paragraph longer form where it fits. Follow all Writing Rulebook rules. Tailor to the JD language and role framing. Label each field answer individually.

OUTPUT: Read each answer. Where it doesn't sound like you, rewrite briefly in your own words. Review individually before entering. Never paste directly without reading first.

2.5 Write cover letter (Amplify — Tech Leading)

FILES: Background · Writing Rulebook · Pipeline · LinkedIn · Advice

PASTE: JD pasted directly. Any company-specific context from networking or research. Direct Claude to include a header, write in paragraphs, and add URLs and contact info at the bottom.

CLAUDE IS TOLD TO: Write a full cover letter drawing from the background and company intelligence provided. Include a header. Paragraph format. Lead with the most relevant framing angle from the JD Analyzer brief. Include all URLs and contact info at the bottom. Follow all Writing Rulebook rules throughout.

OUTPUT: Read full output out loud. Rewrite where sentences don't sound like you. Review full output before sending. Never send without a personal read.

ON REVIEWING ALL OUTPUTS

Every output from every Amplify step, even if technology leading, in this section must be read in full before anything goes anywhere. Read each resume bullet out loud. Read each portal answer out loud. Read the full cover letter out loud. Where something doesn't sound like you, rewrite it in your own words before saving or sending. The system made you faster and more consistent. It did not make you generic. The voice in every output is yours or it doesn't go out.

The Complete Workflow. The Application Process. Part 2.



AUTOMATE



AMPLIFY



PRESERVE

2. THE APPLICATION PROCESS CONTINUED

From identifying contacts through logging the application.

ON NETWORKING AND THE PIPELINE: Every outreach message must feel personal not templated. The LinkedIn Outreach master prompt requires a specific real reference from the contact's actual profile. Generic openers are explicitly banned. The pipeline review and tracker log steps at the end of this section should happen after every single application without exception. The pipeline file is the memory of the entire search. If it is not updated it becomes stale and the system loses its context advantage.

2.6 Identify contacts on LinkedIn (Automate)

FILES: Pipeline · LinkedIn · Advice

PASTE: Raw LinkedIn people list pasted directly. Found by searching the company name alongside talent acquisition, recruiter, the team applied to, the role, the location, and main similar titles, then copying the full list and profile previews LinkedIn provides.

CLAUDE IS TOLD TO: Review the full list and identify the top three contacts to reach out to. Explain reasoning based on title, proximity to the role, and likelihood of being able to help. Rank in priority order.

OUTPUT: Review Claude's contact selections and reasoning before proceeding.

2.7 Select outreach method (Automate)

FILES: Pipeline · Advice

PASTE: What access is currently available: InMail credits, email found, or connection note only. Anecdotally what has been working best recently.

CLAUDE IS TOLD TO: Advise on the best outreach method given availability and what the pipeline history shows has been effective. Give a clear recommendation with brief reasoning.

OUTPUT: Select method and proceed to outreach.

2.8 Write outreach messages (Amplify — Human Leading)

FILES: Writing Rulebook · Pipeline · Background · LinkedIn · GitHub · Advice

PASTE: LinkedIn Outreach master prompt pasted directly. Full LinkedIn profile of the specific contact pasted directly. JD pasted directly.

CLAUDE IS TOLD TO: Draft an InMail or connection note following every rule in the LinkedIn Outreach master prompt. Reference something specific from the contact's profile. Connect most relevant experience to their work. Direct ask matched to seniority. Under 200 words for InMail or 300 chars for connection note. No generic openers. State which type of ask was chosen and why before showing the draft.

OUTPUT: Read full output. Rewrite any sentence that does not sound genuinely like you. Must feel personal not templated. Send the message.

2.9 Full pipeline review and update (Automate)

FILES: Pipeline

PASTE: Pipeline Management master prompt pasted directly. New developments in plain text: emails received, calls completed, rejections, responses, interviews scheduled.

CLAUDE IS TOLD TO: Follow the Pipeline Management master prompt exactly. Update each active company gameplan. Flag any companies crossing from active to dormant or closed. Identify time-sensitive actions needed today. Output the full rewritten gameplan for each changed company.

OUTPUT: Proofread each updated entry. Notes: find each company section, replace exactly. Reupload Pipeline next session.

2.10 Log tracker entry (Automate)

FILES: Pipeline

PASTE: JD pasted directly. Any networking context gathered so far in plain text.

CLAUDE IS TOLD TO: Output one standardized tracker entry in the exact pipeline log format: company, role, date applied, method, contacts reached, status, next action and date.

OUTPUT: Proofread entry. Notes: add to master application log. Reupload Pipeline next session.

ON THE PIPELINE FILE: The pipeline file is the most actively maintained file in the entire system. It updates after every application, every networking message, every response, every interview, every development of any kind. It is what gives the system memory across companies and time. If it falls behind, the system loses its single biggest advantage over a generic AI tool: knowing exactly where things stand with every company at every moment.

The Complete Workflow. Interview Preparation.

 AUTOMATE  AMPLIFY  PRESERVE

3. INTERVIEW PREPARATION

Everything from building the answer bank through the actual interview.

3.1 Company-specific prep (Amplify — Technology Leading)

FILES: Pipeline · Background · AI Experience (for AI-facing roles only) · Advice

PASTE: Interviewer LinkedIn profile pasted directly. Company intelligence gathered in plain text. The specific angle you want to take for this interview in plain text.

CLAUDE IS TOLD TO: Generate a company-specific question bank drawing from the role, the interviewer background, and pipeline history with this company. Generate talking points for each likely question. Request an internet search for typical interview questions in this industry and role to expand the list.

OUTPUT: Review full question bank and talking points. Save company-specific prep to a new note file for reference during drilling. Do not upload back to the project.

3.3 Practice out loud (Preserve — Assisted)

FILES: None automatically loaded for this step.

PASTE: All finalized questions and answers pasted in full. Ask Claude to conduct a mock interview one question at a time. Use the voice recording feature (not voice chat. Voice chat is worse at referencing full chat history) to transcribe your spoken response to each question without a script.

CLAUDE IS TOLD TO: Conduct the mock interview one question at a time. Wait for the transcribed spoken response before continuing. After each response write back specifically where the answer can be improved: what was missing, generic, did not land.

OUTPUT: Upload each transcription. Claude writes back where to improve. Drill repeatedly. Answers refined here are not auto-saved back into the system but can build out more standard question types for future reference.

ON PRACTICING OUT LOUD: The voice recording feature is used instead of voice chat because voice chat is worse at referencing full chat history as specifically as the text interface. Recording your spoken answer separately and uploading the transcription gives Claude the full context of both the prepared answer and what you actually said, which is what makes the feedback useful. The gap between your written answer and your spoken answer is exactly what needs to be closed before the actual interview.

AUTOMATE · AMPLIFY · PRESERVE DEFINED:

- **Automate:** Rule-based and mechanical, output identical regardless of who built the system. Purely for viewing or organization
- **Amplify Technology Leading:** Claude completes process based on extensive ruleset for rule-based task, human reviews and corrects, rewriting always still necessary.
- **Amplify Human Leading:** Human provides raw reasoning, perspective and voice first, Claude structures and refines what the human brings. Human reviews and rewrites where needed.
- **Preserve:** Fully human, Claude assists only in preparation before and debrief after, never touches the moment itself. The key question for every step: is technology doing the work and the human checking it, or is the human doing the work and technology structuring it.

3.2 Build or refine answer bank (Amplify — Human Leading)

FILES: Background · Writing Rulebook · AI Experience · Pipeline

PASTE: JD pasted directly. Ask Claude to generate a full question bank from every responsibility bullet in the JD plus behavioral and industry questions consolidated into one list.

CLAUDE IS TOLD TO: Generate the full question bank first. Then for each question output a complete spoken-voice answer drawn from the full background, written in first person as it would be spoken out loud. Follow all Writing Rulebook rules. Flag any question where the background does not provide strong material.

OUTPUT: Proofread each answer. Read every answer out loud. Where it does not sound like you, rewrite in extreme depth in your own words exactly as you would speak it. Add personality and warmth before saving. Save all Q&A to a dedicated note file.

3.4 The actual interview (Preserve)

FILES: None. Fully human. Preserve tier. Technology does not touch the moment itself.

PASTE (24h after): Notes from the interview in plain text: what was asked, how it went, what landed, what did not, any commitments made, any follow-up promised.

CLAUDE IS TOLD TO (24h after): Rewrite the company pipeline section in full to reflect everything that happened. Update strategic position, current status, and next phases based on the interview outcome.

OUTPUT: Proofread Claude's output. Notes: find that company section, replace it exactly. Reupload the full Pipeline file to Claude next session.

ON THE ACTUAL INTERVIEW: The 24-hour debrief is not optional. The pipeline file needs to reflect what actually happened as soon as possible while details are fresh. Every commitment made in the interview, every question that caught you off guard, every signal about fit or next steps belongs in the pipeline entry immediately. The system cannot help you with follow-up strategy if it does not know what happened.

The Complete Workflow. Sections 4, 5, and General Learning (6).



AUTOMATE



AMPLIFY



PRESERVE

4. WHENEVER NEEDED

Steps used throughout the process as situations arise. Not in fixed chronological order.

4.2 Phone and networking calls (Preserve)

FILES: Pipeline · Background · LinkedIn · AI Experience (AI-facing companies) · Advice

PASTE (before): Person's LinkedIn profile pasted directly. Call context and purpose in plain text. Ask Claude for questions to ask and talking points.

PASTE (after): Call notes in plain text. Ask Claude to rewrite the company pipeline section to reflect the call outcome.

CLAUDE IS TOLD TO: Before: generate specific questions and talking points connecting your background to their work.

After: rewrite the company pipeline section in full to reflect the outcome.

OUTPUT: Proofread pipeline output. Notes: find that company section, replace exactly. Reupload Pipeline next session.

4.4 Offer negotiation (Amplify — Human Leading)

FILES: Pipeline · Background · Advice

PASTE: Exact offer details in plain text including compensation, title, start date, and any other terms.

CLAUDE IS TOLD TO: Advise on whether to negotiate based on the offer details and pipeline context. If negotiating, draft the exact message to send with clear reasoning for the recommendation.

OUTPUT: Review Claude's recommendation carefully. Any negotiation message drafted by Claude must be rewritten meaningfully in your own voice before sending. Never send without a full personal rewrite.

5.1 Update pipeline on any development (Automate)

FILES: Pipeline

PASTE: Pipeline Management master prompt pasted directly. Specific development in plain text: exact email received verbatim, call outcome in full, rejection, interview scheduled, offer received, any status change.

CLAUDE IS TOLD TO: Follow the Pipeline Management master prompt exactly. Rewrite the specific company gameplan section in full. Update current status, strategic position, phases, behavioral rules, and strategic summary. Flag any time-sensitive actions.

OUTPUT: Proofread output. Notes: find that company section, replace it exactly. Reupload the full Pipeline file next session.

4.1 Niche situation advice and strategy (Amplify — Technology Leading)

FILES: Background · AI Experience · Pipeline · LinkedIn · GitHub · Advice

PASTE: Specific situation described in plain text with as much detail as possible. Ask Claude to advise using the relevant files.

CLAUDE IS TOLD TO: Advise on the specific situation drawing from all relevant files. Give a clear recommendation with reasoning. Flag if the advice requires a pipeline update.

OUTPUT: Review Claude's advice carefully. Apply your own judgment before acting on any recommendation. If the advice results in a pipeline update, trigger Section 5.

4.3 Thank you notes and follow-ups (Amplify — Human Leading)

FILES: Writing Rulebook · Pipeline · Background · Advice

PASTE: Specific details from the conversation in plain text so the note is genuinely personalized to what actually happened.

CLAUDE IS TOLD TO: Draft a thank you note drawing from the specific details provided. Follow all Writing Rulebook rules. Reference something specific from the actual conversation. No generic closers.

OUTPUT: Read full output. Rewrite any sentence that does not sound genuinely like you before sending. Log send date in Notes. Find that company section, add send date and status update. Reupload Pipeline next session.

5. WHENEVER ANY UPDATE HAPPENS

Triggered any time a new development occurs across any active company regardless of where in the process you are.

6.1 General learning workflow (Amplify — Human Leading)

ABOUT: This is how I approach learning alongside AI rather than having it do the thinking for me as I gain new experience and skills. Section 6: Claude can be used as a learning partner for any subject relevant to your work. Have Claude test your reasoning rather than just produce answers.

FILES: AI Experience · Background · Advice

PASTE: How to Use Claude master prompt pasted directly. Specific concept or skill you are trying to learn in plain text. Ask Claude to quiz your reasoning rather than just explain.

CLAUDE IS TOLD TO: Quiz reasoning rather than explain concepts. Give problems to work through. Push back on answers. Identify where reasoning breaks down before confirming or correcting. Never give the answer before the person has attempted it.

OUTPUT: Ask Claude to identify gaps in your understanding before giving the answer. Accelerate understanding, not outsource the thinking itself.

The Results. The Honest Picture.

THE EFFICIENCY REALITY

23

copy pastes maximum per company for the entire job obtaining process



6-15

steps for applications alone



100%

of all outputs reviewed before anything goes out

Every copy paste produces an output that gets read, checked, and revised before anything goes anywhere. You are not blindly submitting AI output. You are vetting every response and rewriting where it does not sound right. The system never submits anything. The person does. It just makes sure you never start from a blank page.

WHAT IT ACTUALLY PRODUCED

WHAT THE SYSTEM PRODUCED

Significantly higher application to interview conversion rate than most people in the same recruiting cycle, starting from a late entry into an already-closed cycle.

Multiple second rounds and live interviews produced across companies in financial analytics, consulting, AI, and sports analytics.

A 30-slide live work showcase presented to a consulting firm's AI team using the firm itself as the case study. Built entirely through the system.

The Automate · Amplify · Preserve framework arrived at independently through building and failing, then validated by encountering the same concepts published by professionals in the field.

A system that can be rebuilt for any person, any industry, any role. The architecture is the insight, not the specific content.

THE HONEST PICTURE

The system gets you in the room. What happens in the room is fully on you. No system changes that.

The first version failed completely. Scrapping it and diagnosing why taught more about AI deployment than any course could have.

The system is only as good as the quality of the context files. Stale files mean stale output. Maintenance is a feature, not a burden.

The hardest part was not the tools. It was building something that actually knows you.

Interview drilling is the next frontier. The system produces excellent preparation material. Using it consistently and deeply is the part still being built.

THE SYSTEM NEVER SUBMITS ANYTHING. I DO.

Read every output

Nothing goes out without being read in full. If it does not sound right it gets rewritten. Not skimmed. Read.

Read bullets out loud

The single most reliable check for whether something sounds like you rather than AI. If you would not say it out loud it does not go out.

Rewrite where needed

Not where it is wrong. Where it does not sound like you. The distinction matters. Generic but accurate is not good enough.

Then submit

The application reflects genuine interest, real fit assessment, and your actual voice. The system made you faster and more consistent. It did not make you generic.

The Philosophy.

Four principles arrived at through building, failing, and rebuilding. Not through reading.

“The quality of AI output never exceeds the quality of the context it is working from.”

Generic inputs produce generic outputs. Building the context architecture, what the system knows, how it is structured, how it is maintained — is the hard part. The tools are the easy part. This is not a prompt engineering insight. It is an architectural one.

“AI systems built for specific people need to know the person in detail.”

Generic architecture produces generic output regardless of how sophisticated the pipeline is. The system does not replace judgment. It gives you your judgment faster. That distinction is everything.

“Specificity is the entire skill.”

Ask Claude to ask questions before it produces anything. Pressure test every output. Push back when something feels generic or off. Be deliberately redundant on rules. Treat it as a fast knowledgeable thinking partner that still needs human judgment to be useful.

“Where AI belongs and where it does not is a design decision, not an afterthought.”

Automating formatting is safe. Automating relationship building is not. Automating intake is safe. Automating judgment calls about fit and voice is not. Every step in the workflow reflects a judgment about what AI can be trusted to do and what it cannot. That judgment, applied systematically, is the system.

AUTOMATE THE INTAKE.
PRESERVE THE THINKING.

The system automates the mechanical so judgment is never wasted on intake.

The system amplifies thinking so every output is better than what I could produce alone.

The system preserves the human moments because no tool should touch what actually matters.

I built this to get here. I want to build the next version of something like this professionally.

[illegible]

Graham Friedman · github.com/grahamdfriedman12 · linkedin.com/in/friedg

This presentation was designed with assistance from Claude in PowerPoint ;)